

```
In [ ]: import pandas as pd
import numpy as np
import torch.nn as nn
import matplotlib.pyplot as plt
import re
import warnings
import html
import tiktoken
import time
import torch
import copy
import math
import time
import pickle
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from torch.utils.data import Dataset, DataLoader
from bs4 import BeautifulSoup, MarkupResemblesLocatorWarning
from ftfy import fix_text

warnings.filterwarnings("ignore", category=MarkupResemblesLocatorWarning)
warnings.simplefilter(action="ignore", category=pd.errors.SettingWithCopyWarning)

torch.manual_seed(123)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
In [ ]: class PhishingEmailPreprocessor:
    """
    Veri ön işleme için yazılan sınıf.
    """

    def __init__(self, csv_path:str, train_frac:float, validation_frac:float):

        self.df = None
        self.train_df = None
        self.validation_df = None
        self.test_df = None

        self.train_frac = train_frac
        self.validation_frac = validation_frac

        self.csv_path = csv_path

        self.tokenizer = tiktoken.get_encoding("gpt2")
        self.pad_token_id = 50256
        self.max_len = 1024

        self.load_and_prepare(self.csv_path)

    def strip_html(self, text: str) -> str:
        if not isinstance(text, str):
            return ""
        text = html.unescape(text)
        soup = BeautifulSoup(text, "html.parser")
        return soup.get_text(separator=" ")
```

```

def normalize_whitespace(self, text: str) -> str:
    if not isinstance(text, str):
        return ""
    text = re.sub(r"\s+", " ", text)
    return text.strip()

def fix_unicode(self, text: str) -> str:
    if not isinstance(text, str):
        return ""
    return fix_text(text)

def create_balanced_dataset(self):

    num_spam = self.df[self.df["Email Type"] == 1].shape[0]

    ham_subset = self.df[self.df["Email Type"] == 0].sample(num_spam, random

    self.df = pd.concat([ham_subset, self.df[self.df["Email Type"] == 1]])

def load_and_prepare(self, csv_path: str):
    self.df = pd.read_csv(csv_path, sep=",", header=0, quotechar='"')

    self.df = self.df.dropna(subset=["Email Text", "Email Type"])
    self.df = self.df[self.df["Email Text"].astype(str).str.strip() != ""]

    self.df["Email Text"] = self.df["Email Text"].apply(self.fix_unicode)
    self.df["Email Text"] = self.df["Email Text"].apply(self.strip_html)
    self.df["Email Text"] = self.df["Email Text"].apply(self.normalize_white

    self.df = self.df[self.df["Email Text"].str.len() > 0]

    self.df = self.df.drop_duplicates(subset=["Email Text"]).copy()

    self.df["Email Type"] = self.df["Email Type"].map({"Safe Email": 0, "Phi
    self.df = self.df.dropna(subset=["Email Type"])
    self.df["Email Type"] = self.df["Email Type"].astype(int)
    self.create_balanced_dataset()
    self.random_split()
    self.train_df["Token"] = self.train_df["Email Text"].apply(self.tokenize
    self.validation_df["Token"] = self.validation_df["Email Text"].apply(sel
    self.test_df["Token"] = self.test_df["Email Text"].apply(self.tokenize)

def random_split(self):
    self.df = self.df.sample(frac=1, random_state=123).reset_index(drop=True

    train_end = int(len(self.df) * self.train_frac)
    validation_end = train_end + int(len(self.df) * self.validation_frac)

    self.train_df = self.df[:train_end]
    self.validation_df = self.df[train_end:validation_end]
    self.test_df = self.df[validation_end:]

def tokenize(self, text):
    tokens = self.tokenizer.encode(text)
    return tokens

def get_dfs(self) -> pd.DataFrame:
    self.train_df = self.train_df.drop(columns=["Unnamed: 0"])
    self.validation_df = self.validation_df.drop(columns=["Unnamed: 0"])

```

```
self.test_df = self.test_df.drop(columns=["Unnamed: 0"])
return self.train_df, self.validation_df, self.test_df
```

```
In [ ]: preprocessor = PhishingEmailPreprocessor(csv_path="Phishing_Email.csv", train_fr
train_df, validation_df, test_df = preprocessor.get_dfs())
```

```
In [ ]: """
Dataframeden DataLoader'a dönüşüm için aşağıdaki 2 sınıf kullanılmıştır.
"""

class EmailDataset(Dataset):
    def __init__(self, df):
        self.df = df.reset_index(drop=True)

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        tokens = self.df.loc[idx, "Token"]
        label = int(self.df.loc[idx, "Email Type"])
        return tokens, label

class EmailDataLoader:
    def __init__(self, train_df, val_df, test_df, batch_size=8, num_workers=0, s
        self.batch_size = batch_size
        self.num_workers = num_workers

        self.train_dataset = EmailDataset(train_df)
        self.val_dataset = EmailDataset(val_df)
        self.test_dataset = EmailDataset(test_df)

    #Bir batchdeki maksimum token uzunluğunu bulup batch padding yapmak için yaz
    def make_collate_fn(self, pad_id=preprocessor.pad_token_id, max_len=None):
        def collate(batch):
            xs, ys = zip(*batch)

            xs_list = [x.tolist() if isinstance(x, torch.Tensor) else x for x in

            if max_len is not None:
                xs_list = [x[:max_len] for x in xs_list]

            maxlen = max(len(x) for x in xs_list)

            X = [x + [pad_id] * (maxlen - len(x)) for x in xs_list]
            Y = list(ys)

            X = torch.tensor(X, dtype=torch.long)
            Y = torch.tensor(Y, dtype=torch.long)
            return X, Y
        return collate

    def get_loaders(self):
        collate = self.make_collate_fn(pad_id=preprocessor.pad_token_id, max_len

        train_loader = DataLoader(
            dataset=self.train_dataset,
            batch_size=self.batch_size,
            shuffle=True,
```

```

        num_workers=self.num_workers,
        drop_last=True,
        collate_fn=collate,
    )

    val_loader = DataLoader(
        dataset=self.val_dataset,
        batch_size=self.batch_size,
        shuffle=False,
        num_workers=self.num_workers,
        drop_last=False,
        collate_fn=collate,
    )

    test_loader = DataLoader(
        dataset=self.test_dataset,
        batch_size=self.batch_size,
        shuffle=False,
        num_workers=self.num_workers,
        drop_last=False,
        collate_fn=collate,
    )

    return train_loader, val_loader, test_loader

```

```
In [ ]: data_module = EmailDataLoader(train_df, validation_df, test_df, batch_size=16)
train_loader, val_loader, test_loader = data_module.get_loaders()
```

```
In [ ]: class MultiHeadAttention(nn.Module):
    """
    Transformer bloğundaki attention.
    """
    def __init__(self, d_in, d_out, context_length, dropout, num_heads, qkv_bias):
        super().__init__()
        assert (d_out % num_heads == 0)

        self.d_out = d_out
        self.num_heads = num_heads
        self.head_dim = d_out // num_heads

        self.W_query = nn.Linear(d_in, d_out, bias=qkv_bias)
        self.W_key = nn.Linear(d_in, d_out, bias=qkv_bias)
        self.W_value = nn.Linear(d_in, d_out, bias=qkv_bias)
        self.out_proj = nn.Linear(d_out, d_out)
        self.dropout = nn.Dropout(dropout)
        self.register_buffer(
            "mask",
            torch.triu(torch.ones(context_length, context_length),
                        diagonal=1)
        )

    def forward(self, x):
        b, num_tokens, d_in = x.shape

        keys = self.W_key(x)
        queries = self.W_query(x)
        values = self.W_value(x)

```

```

keys = keys.view(b, num_tokens, self.num_heads, self.head_dim)
values = values.view(b, num_tokens, self.num_heads, self.head_dim)
queries = queries.view(b, num_tokens, self.num_heads, self.head_dim)

keys = keys.transpose(1, 2)
queries = queries.transpose(1, 2)
values = values.transpose(1, 2)

attn_scores = queries @ keys.transpose(2, 3)

mask_bool = self.mask.bool()[ :num_tokens, :num_tokens]

attn_scores.masked_fill_(mask_bool, -torch.inf)

attn_weights = torch.softmax(attn_scores / keys.shape[-1]**0.5, dim=-1)
attn_weights = self.dropout(attn_weights)

context_vec = (attn_weights @ values).transpose(1, 2)

context_vec = context_vec.contiguous().view(b, num_tokens, self.d_out)
context_vec = self.out_proj(context_vec) # optional projection

return context_vec

```

```

In [ ]: class GELU(nn.Module):
    """
    Aktivasyon fonksiyonu için yazılmış sınıf.
    """
    def __init__(self):
        super().__init__()

    def forward(self, x):
        return 0.5 * x * (1 + torch.tanh(
            torch.sqrt(torch.tensor(2.0 / torch.pi)) *
            (x + 0.044715 * torch.pow(x, 3))
        ))

class FeedForward(nn.Module):
    def __init__(self, cfg):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Linear(cfg["emb_dim"], 4 * cfg["emb_dim"]),
            GELU(),
            nn.Linear(4 * cfg["emb_dim"], cfg["emb_dim"]),
        )

    def forward(self, x):
        return self.layers(x)

class LayerNorm(nn.Module):
    def __init__(self, emb_dim):
        super().__init__()
        self.eps = 1e-5
        self.scale = nn.Parameter(torch.ones(emb_dim))
        self.shift = nn.Parameter(torch.zeros(emb_dim))

```

```

def forward(self, x):
    mean = x.mean(dim=-1, keepdim=True)
    var = x.var(dim=-1, keepdim=True, unbiased=False)
    norm_x = (x - mean) / torch.sqrt(var + self.eps)
    return self.scale * norm_x + self.shift

class TransformerBlock(nn.Module):
    def __init__(self, cfg):
        super().__init__()
        self.att = MultiHeadAttention(
            d_in=cfg["emb_dim"],
            d_out=cfg["emb_dim"],
            context_length=cfg["context_length"],
            num_heads=cfg["n_heads"],
            dropout=cfg["drop_rate"],
            qkv_bias=cfg["qkv_bias"])
        self.ff = FeedForward(cfg)
        self.norm1 = LayerNorm(cfg["emb_dim"])
        self.norm2 = LayerNorm(cfg["emb_dim"])
        self.drop_shortcut = nn.Dropout(cfg["drop_rate"])

    def forward(self, x):
        shortcut = x
        x = self.norm1(x)
        x = self.att(x)
        x = self.drop_shortcut(x)
        x = x + shortcut

        shortcut = x
        x = self.norm2(x)
        x = self.ff(x)
        x = self.drop_shortcut(x)
        x = x + shortcut

        return x

class GPTModel(nn.Module):
    """
    Kullanılacak model.
    """
    def __init__(self, cfg):
        super().__init__()
        self.tok_emb = nn.Embedding(cfg["vocab_size"], cfg["emb_dim"])
        self.pos_emb = nn.Embedding(cfg["context_length"], cfg["emb_dim"])
        self.drop_emb = nn.Dropout(cfg["drop_rate"])

        self.trf_blocks = nn.Sequential(
            *[TransformerBlock(cfg) for _ in range(cfg["n_layers"])]])

        self.final_norm = LayerNorm(cfg["emb_dim"])
        self.out_head = nn.Linear(
            cfg["emb_dim"], cfg["vocab_size"], bias=False
        )

    def forward(self, in_idx):
        batch_size, seq_len = in_idx.shape
        tok_embeddings = self.tok_emb(in_idx)
        pos_embeddings = self.pos_emb(torch.arange(seq_len, device=in_idx.device))
        x = tok_embeddings + pos_embeddings # Shape [batch_size, num_tokens, emb_size]
        x = self.drop_emb(x)

```

```

        x = self.trf_blocks(x)
        x = self.final_norm(x)
        logits = self.out_head(x)
        return logits

```

```

In [ ]: #Modeli huggingface'den indirip config bilgisi ile yükledim.
file_name = "gpt2-small-124M.pth"
GPT_CONFIG_124M = {
    "vocab_size": 50257, # Vocabulary size
    "context_length": 1024, # Shortened context length (orig: 1024)
    "emb_dim": 768, # Embedding dimension
    "n_heads": 12, # Number of attention heads
    "n_layers": 12, # Number of layers
    "drop_rate": 0.1, # Dropout rate
    "qkv_bias": True # Query-key-value bias
}

model = GPTModel(GPT_CONFIG_124M)
model.load_state_dict(torch.load(file_name, weights_only=True))
model.eval()
model.to(device)

```

```

In [ ]: #Eğitime başlamadan önce tüm ağırlıkları dondurdum.
for param in model.parameters():
    param.requires_grad = False

```

```

In [ ]: #Modelin son kısmını 2 çıkışlı network ile değiştirdim.
num_classes = 2
model.out_head = torch.nn.Linear(in_features=GPT_CONFIG_124M["emb_dim"], out_features=num_classes)
model.out_head = model.out_head.to(device)

```

```

In [ ]: #Son transformer block ve finalLayerNorm kısmını eğitim amaçlı açtım.
for param in model.trf_blocks[-1].parameters():
    param.requires_grad = True

for param in model.final_norm.parameters():
    param.requires_grad = True

```

```

In [ ]: PAD_ID = 50256

#Modelde attentionda PAD mask olmadığı için en son valid token almak için yazıldı
def last_valid_logits(xb, model, device, pad_id=PAD_ID):
    xb = xb.to(device)
    last_idx = (xb != pad_id).sum(dim=-1) - 1
    logits_all = model(xb) # [B, T, C]
    return logits_all[tuple(torch.arange(xb.size(0), device=device), last_idx, :)] # [B, C]

#cross_entropy loss kullandım.
def calc_loss_batch(input_batch, target_batch, model, device):
    input_batch, target_batch = input_batch.to(device), target_batch.to(device)
    logits = last_valid_logits(input_batch, model, device, PAD_ID)
    loss = torch.nn.functional.cross_entropy(logits, target_batch)
    return loss

def calc_loss_loader(data_loader, model, device, num_batches=None):
    total_loss = 0.
    if len(data_loader) == 0:

```

```

        return float("nan")
    elif num_batches is None:
        num_batches = len(data_loader)
    else:
        num_batches = min(num_batches, len(data_loader))
    for i, (input_batch, target_batch) in enumerate(data_loader):
        if i < num_batches:
            loss = calc_loss_batch(input_batch, target_batch, model, device)
            total_loss += loss.item()
        else:
            break
    return total_loss / num_batches

def calc_accuracy_loader(data_loader, model, device, num_batches=None):
    model.eval()
    correct_predictions, num_examples = 0, 0

    if num_batches is None:
        num_batches = len(data_loader)
    else:
        num_batches = min(num_batches, len(data_loader))
    for i, (input_batch, target_batch) in enumerate(data_loader):
        if i < num_batches:
            input_batch, target_batch = input_batch.to(device), target_batch.to(device)

            with torch.no_grad():
                logits = last_valid_logits(input_batch, model, device, PAD_ID)
                predicted_labels = torch.argmax(logits, dim=-1)

            num_examples += predicted_labels.shape[0]
            correct_predictions += (predicted_labels == target_batch).sum().item()
        else:
            break
    return correct_predictions / num_examples

def evaluate_model(model, train_loader, val_loader, device, eval_iter):
    model.eval()
    with torch.no_grad():
        train_loss = calc_loss_loader(train_loader, model, device, num_batches=eval_iter)
        val_loss = calc_loss_loader(val_loader, model, device, num_batches=eval_iter)
    model.train()
    return train_loss, val_loss

```

In []: *#Eğitim için early stopping ile kullanacağımız fonksiyon*

```

def train_classifier_simplev2(
    model, train_loader, val_loader, optimizer, device,
    num_epochs, eval_freq, eval_iter,
    early_stopping=True, patience=3, min_delta=1e-3,
    restore_best_weights=True, checkpoint_path=None
):
    # Takip Listeleri
    train_losses, val_losses, train_accs, val_accs = [], [], [], []
    examples_seen, global_step = 0, -1

    # Early stopping durum değişkenleri
    best_val_loss = math.inf
    best_state_dict = None
    epochs_no_improve = 0
    best_epoch = -1

```



```

for epoch in range(num_epochs):
    model.train()

    for input_batch, target_batch in train_loader:
        optimizer.zero_grad()
        loss = calc_loss_batch(input_batch, target_batch, model, device)
        loss.backward()
        optimizer.step()
        examples_seen += input_batch.shape[0]
        global_step += 1

    if global_step % eval_freq == 0 and global_step > 0:

        train_loss, val_loss = evaluate_model(
            model, train_loader, val_loader, device, eval_iter=eval_iter
        )
        train_losses.append(train_loss)
        val_losses.append(val_loss)
        print(f"Ep {epoch+1} (Step {global_step:06d}): "
              f"Train loss {train_loss:.3f}, Val loss {val_loss:.3f}")

    epoch_train_loss, epoch_val_loss = evaluate_model(
        model, train_loader, val_loader, device, eval_iter=None # tamamını
    )

    train_accuracy = calc_accuracy_loader(train_loader, model, device, num_b
    val_accuracy = calc_accuracy_loader(val_loader, model, device, num_b
    print(f"[Epoch {epoch+1}] mean Train loss {epoch_train_loss:.3f}, mean V
    print(f"Training accuracy: {train_accuracy*100:.2f}% | Validation accura

    train_accs.append(train_accuracy)
    val_accs.append(val_accuracy)

    # ---- EARLY STOPPING KONTROLÜ ----
    if early_stopping:
        improved = (best_val_loss - epoch_val_loss) > min_delta
        if improved:
            best_val_loss = epoch_val_loss
            best_epoch = epoch + 1
            epochs_no_improve = 0
            if restore_best_weights:
                best_state_dict = copy.deepcopy(model.state_dict())
            if checkpoint_path: # opsiyonel diske kaydet
                torch.save(model.state_dict(), checkpoint_path)
        else:
            epochs_no_improve += 1
            if epochs_no_improve >= patience:
                print(f"Early stopping: {patience} epoch boyunca anlamlı iyi
                      f"En iyi Val loss {best_val_loss:.4f} (epoch {best_epo
                if restore_best_weights and best_state_dict is not None:
                    model.load_state_dict(best_state_dict)
                    print("Best weights geri yüklendi.")
                break

    # ---- /EARLY STOPPING ----

    return train_losses, val_losses, train_accs, val_accs, examples_seen

```

```
In [ ]: start_time = time.time()

#Optimizer olarak AdamW kullandım ve lr ile weight_decay değerlerini derste gördüm
optimizer = torch.optim.AdamW(model.parameters(), lr=1e-5, weight_decay=0.01)

#ilk eğitimde 5 aldım. 5.epocha kadar Validation Loss düştüğünü gördüm networkün
num_epochs = 20
train_losses, val_losses, train_accs, val_accs, examples_seen = train_classifier(
    model, train_loader, val_loader, optimizer, device,
    num_epochs=num_epochs,
    eval_freq=200,      # step içinde kısa örnekleme rapor
    eval_iter=4,        # step içi değerlendirme kaç batch taransın
    early_stopping=True,
    patience=3,          # 3 epoch iyileşme yoksa dur
    min_delta=1e-3,      # val loss en az bu kadar düşerse "iyileşme" say
    restore_best_weights=True,
    checkpoint_path="phishing_classifier2.pth" # istersen "best.pt" ver
)

end_time = time.time()
execution_time_minutes = (end_time - start_time) / 60
print(f"Training completed in {execution_time_minutes:.2f} minutes.")
```

```
In [ ]: #Lazım olabilir.
torch.save(model.state_dict(), "phishing_classifier2.pth")
metrics = { "train_losses": train_losses, "val_losses": val_losses, "train_accs":
with open("metrics2.pkl", "wb") as f:
    pickle.dump(metrics, f, protocol=pickle.HIGHEST_PROTOCOL)
```

```
In [ ]: @torch.no_grad()
def eval_full(val_loader, model, device, pad_id=50256):
    model.eval()
    tp = fp = tn = fn = 0
    all_probs = []
    all_y = []

    for xb, yb in val_loader:
        xb = xb.to(device, non_blocking=True)
        yb = yb.to(device, non_blocking=True)

        # Son geçerli token'ın indeksini bul
        last_idx = (xb != pad_id).sum(dim=1) - 1

        # Logits'i o pozisyonlardan al
        logits_all = model(xb) # [B, T, C]
        logits = logits_all[torch.arange(xb.size(0), device=device), last_idx, :]

        preds = torch.argmax(logits, dim=-1) # [B]
        probs = torch.softmax(logits, dim=-1)[: , 1] # positive sınıf

        # Confusion matrix bileşenleri
        tp += ((preds == 1) & (yb == 1)).sum().item()
        tn += ((preds == 0) & (yb == 0)).sum().item()
        fp += ((preds == 1) & (yb == 0)).sum().item()
        fn += ((preds == 0) & (yb == 1)).sum().item()

    all_probs.append(probs.detach().cpu())
    all_y.append(yb.detach().cpu())
```

```

# Temel metrikler
total = tp + tn + fp + fn
acc = (tp + tn) / max(1, total)
precision = tp / max(1, tp + fp)
recall = tp / max(1, tp + fn)
f1 = 2 * precision * recall / max(1e-12, (precision + recall))

# ROC-AUC (ranks yöntemi; sklearn yok)
y_true = torch.cat(all_y).numpy()
scores = torch.cat(all_probs).numpy()
pos = (y_true == 1)
n_pos = int(pos.sum())
n_neg = len(y_true) - n_pos

auc = None
if n_pos > 0 and n_neg > 0:
    order = np.argsort(scores)
    ranks = np.empty_like(order)
    ranks[order] = np.arange(len(scores)) + 1
    sum_ranks_pos = ranks[pos].sum()
    auc = (sum_ranks_pos - n_pos * (n_pos + 1) / 2) / (n_pos * n_neg)

return {
    "accuracy": acc,
    "precision": precision,
    "recall": recall,
    "f1": f1,
    "roc_auc": auc,
    "confusion_matrix": {"tn": tn, "fp": fp, "fn": fn, "tp": tp},
}

```

```

In [ ]: validation_metrics = eval_full(val_loader, model, device)
        test_metrics = eval_full(test_loader, model, device)

```

```

In [ ]: #TF-IDF vs fine-tuned GPT2-small
X_train = train_df["Email Text"].astype(str).values
y_train = train_df["Email Type"].astype(int).values

X_val    = validation_df["Email Text"].astype(str).values
y_val    = validation_df["Email Type"].astype(int).values

X_test   = test_df["Email Text"].astype(str).values
y_test   = test_df["Email Type"].astype(int).values

pipe = Pipeline([
    ("tfidf", TfidfVectorizer(
        lowercase=False,
        strip_accents="unicode",
        sublinear_tf=True
    )),
    ("lr", LogisticRegression(
        max_iter=2000,
        n_jobs=None,
        solver="liblinear"
    ))
])

param_grid = {

```

```

        "tfidf__ngram_range": [(1,1), (1,2)],
        "tfidf__min_df": [2, 5],
        "tfidf__max_features": [None, 100_000],
        "lr__C": [0.5, 1.0, 2.0],
        "lr__penalty": ["l2"]
    }

gs = GridSearchCV(
    pipe,
    param_grid=param_grid,
    scoring="f1",
    cv=3,
    n_jobs=-1,
    verbose=1
)
gs.fit(X_train, y_train)

baseline = gs.best_estimator_

# === 3) Validation ve Test performansı ===
def evaluate(model, X, y):
    proba = model.predict_proba(X)[:,-1]
    pred = (proba >= 0.5).astype(int)
    return {
        "accuracy": accuracy_score(y, pred),
        "precision": precision_score(y, pred, zero_division=0),
        "recall": recall_score(y, pred, zero_division=0),
        "f1": f1_score(y, pred, zero_division=0),
        "roc_auc": roc_auc_score(y, proba),
        "confusion_matrix": confusion_matrix(y, pred, labels=[0,1])
    }

baseline_val = evaluate(baseline, X_val, y_val)
baseline_test = evaluate(baseline, X_test, y_test)

# GPT2 test veri setindeki sonuçlarım.
gpt2_test = {
    "accuracy": 0.9686544342507645,
    "precision": 0.9734375,
    "recall": 0.9629057187017002,
    "f1": 0.9681429681429681,
    "roc_auc": 0.9965627462488338,
    "confusion_matrix": np.array([[644, 17],[24, 623]])
}

def plot_compare(baseline_metrics, gpt2_metrics, title_suffix="(Test Set)":
    labels = ["Accuracy", "Precision", "Recall", "F1", "ROC-AUC"]
    base_vals = [baseline_metrics[k.lower()] for k in ["ACCURACY", "PRECISION", "RECALL", "F1", "ROC-AUC"]]
    gpt2_vals = [gpt2_metrics[k.lower()] for k in ["ACCURACY", "PRECISION", "RECALL", "F1", "ROC-AUC"]]

    x = np.arange(len(labels))
    w = 0.35
    plt.figure(figsize=(9,5))
    b1 = plt.bar(x - w/2, base_vals, width=w, label="TF-IDF + LR")
    b2 = plt.bar(x + w/2, gpt2_vals, width=w, label="GPT-2 Fine-tuned")
    for bars in [b1, b2]:
        for b in bars:
            h = b.get_height()
            plt.text(b.get_x() + b.get_width()/2, h+0.01, f"{h:.2f}", ha="center")

```

```
plt.ylim(0,1.05)
plt.xticks(x, labels)
plt.title(f"Performans Karşılaştırması {title_suffix}")
plt.ylabel("Skor")
plt.legend()
plt.grid(axis="y", alpha=0.2)
plt.tight_layout()
plt.show()

plot_compare(baseline_test, gpt2_test, "(Test)")
```